

Vegetable image classifier using MobileNetV3

```
In [1]: #import data by kaggle  
!mkdir -p ~/.kaggle  
!cp kaggle.json ~/.kaggle/
```

```
In [2]: !kaggle datasets download -d misrakahmed/vegetable-image-dataset  
  
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'  
Downloading vegetable-image-dataset.zip to /content  
97% 519M/534M [00:05<00:00, 69.4MB/s]  
100% 534M/534M [00:05<00:00, 101MB/s]
```

```
In [3]: #file unzip  
import zipfile  
zip_ref = zipfile.ZipFile('/content/vegetable-image-dataset.zip')  
zip_ref.extractall('/content')  
zip_ref.close()
```

```
In [7]: import matplotlib.pyplot as plt  
import numpy as np  
import pandas as pd  
from sklearn.model_selection import train_test_split  
import os  
import PIL  
import pathlib  
import tensorflow as tf  
from tensorflow import keras
```

```
In [8]: train_dir = pathlib.Path('/content/Vegetable Images/train')  
val_dir = pathlib.Path('/content/Vegetable Images/validation')  
test_dir = pathlib.Path('/content/Vegetable Images/test')  
train_image_count = len(list(train_dir.glob('*/*.jpg')))  
train_image_count
```

```
Out[8]: 15000
```

```
In [9]: batch_size = 32
image_width = 224
image_height = 224

train_ds = keras.preprocessing.image_dataset_from_directory( train_
    seed=123,
    #validation_split=0.5, #subset='training',
    image_size=(image_height, image_width),
    batch_size=batch_size
)
val_ds = keras.preprocessing.image_dataset_from_directory( val_dir,
    seed=123,
    image_size=(image_height, image_width),
    batch_size=batch_size
)

test_ds = keras.preprocessing.image_dataset_from_directory( test_di
    seed=123,
    image_size=(image_height, image_width),
    batch_size=batch_size
)
```

```
Found 15000 files belonging to 15 classes.
Found 3000 files belonging to 15 classes.
Found 3000 files belonging to 15 classes.
```

```
In [10]: #check data class
class_names = train_ds.class_names
class_names
```

```
Out[10]: ['Bean',
'Bitter_Gourd',
'Bottle_Gourd',
'Brinjal',
'Broccoli',
'Cabbage',
'Capsicum',
'Carrot',
'Cauliflower',
'Cucumber',
'Papaya',
'Potato',
'Pumpkin',
'Radish',
'Tomato']
```

```
In [11]: import matplotlib.pyplot as plt
plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i+1)
        plt.imshow(images[i].numpy().astype('uint8'))
        plt.title(class_names[labels[i]])
        plt.axis('off')
```

Broccoli



Cabbage



Tomato



Broccoli



Bean



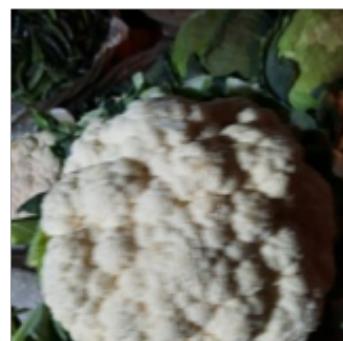
Broccoli



Bean



Cauliflower



Carrot



```
In [12]: # Criando o modelo base em cima do modelo MobileNetV3
base_model = keras.applications.MobileNetV3Small(input_shape=(image
    classes=400,
    include_top=False,
    weights='imagenet')
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v3/weights_mobilenet_v3_small_224_1.0_float_no_top_v2.h5
(https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v3/weights_mobilenet_v3_small_224_1.0_float_no_top_v2.h5)
4334752/4334752 [=====] - 0s 0us/step

```
In [13]: # Freeze convolutional base
base_model.trainable = False
base_model.summary()
```

Model: "MobilenetV3small"

Layer (type) connected to	Output Shape	Param #	C
input_1 (InputLayer)	[(None, 224, 224, 3 0)]		[
rescaling (Rescaling) 'input_1[0][0]'	(None, 224, 224, 3) 0		[
Conv (Conv2D) 'rescaling[0][0]'	(None, 112, 112, 16 432)		[
Conv/BatchNorm (BatchNormaliza 'conv[0][0]'	(None, 112, 112, 16 64)		[

```
In [14]: data_augmentation = keras.models.Sequential([
    keras.layers.RandomFlip('horizontal'),
    keras.layers.RandomRotation(0.2)
])
```

```
In [15]: num_classes = len(class_names)# 7

inputs = keras.Input(shape=(image_width, image_height, 3))
#x = data_augmentation(inputs)
x = keras.applications.mobilenet_v3.preprocess_input(inputs)
x = base_model(x, training=False)
x = keras.layers.GlobalAveragePooling2D()(x)
x = keras.layers.Dropout(0.2)(x)

outputs = keras.layers.Dense(num_classes, activation='softmax')(x)
model = keras.Model(inputs, outputs)
```

```
In [16]: #compile model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
MobilenetV3small (Functiona l)	(None, 7, 7, 576)	939120
global_average_pooling2d (G lobalAveragePooling2D)	(None, 576)	0
dropout (Dropout)	(None, 576)	0
dense (Dense)	(None, 15)	8655
<hr/>		
Total params: 947,775		
Trainable params: 8,655		
Non-trainable params: 939,120		

```
In [17]: #fitting model
initial_epochs = 10

early_stop = keras.callbacks.EarlyStopping(patience=1, restore_best_weights=True)

history = model.fit(train_ds,
                     validation_data=val_ds,
                     epochs=initial_epochs,
                     callbacks=[early_stop])
```

```
Epoch 1/10
469/469 [=====] - 32s 46ms/step - loss: 0.4650 - accuracy: 0.8921 - val_loss: 0.0958 - val_accuracy: 0.9850
Epoch 2/10
469/469 [=====] - 21s 44ms/step - loss: 0.0846 - accuracy: 0.9876 - val_loss: 0.0458 - val_accuracy: 0.9930
Epoch 3/10
469/469 [=====] - 21s 44ms/step - loss: 0.0500 - accuracy: 0.9924 - val_loss: 0.0305 - val_accuracy: 0.9947
Epoch 4/10
469/469 [=====] - 21s 44ms/step - loss: 0.0345 - accuracy: 0.9945 - val_loss: 0.0225 - val_accuracy: 0.9960
Epoch 5/10
469/469 [=====] - 21s 44ms/step - loss: 0.0275 - accuracy: 0.9951 - val_loss: 0.0181 - val_accuracy: 0.9963
Epoch 6/10
469/469 [=====] - 20s 42ms/step - loss: 0.0211 - accuracy: 0.9970 - val_loss: 0.0156 - val_accuracy: 0.9963
Epoch 7/10
469/469 [=====] - 21s 44ms/step - loss: 0.0167 - accuracy: 0.9979 - val_loss: 0.0131 - val_accuracy: 0.9973
Epoch 8/10
469/469 [=====] - 20s 42ms/step - loss: 0.0139 - accuracy: 0.9979 - val_loss: 0.0106 - val_accuracy: 0.9977
Epoch 9/10
469/469 [=====] - 21s 44ms/step - loss: 0.0123 - accuracy: 0.9978 - val_loss: 0.0106 - val_accuracy: 0.9983
Epoch 10/10
469/469 [=====] - 20s 42ms/step - loss: 0.0107 - accuracy: 0.9984 - val_loss: 0.0096 - val_accuracy: 0.9977
```

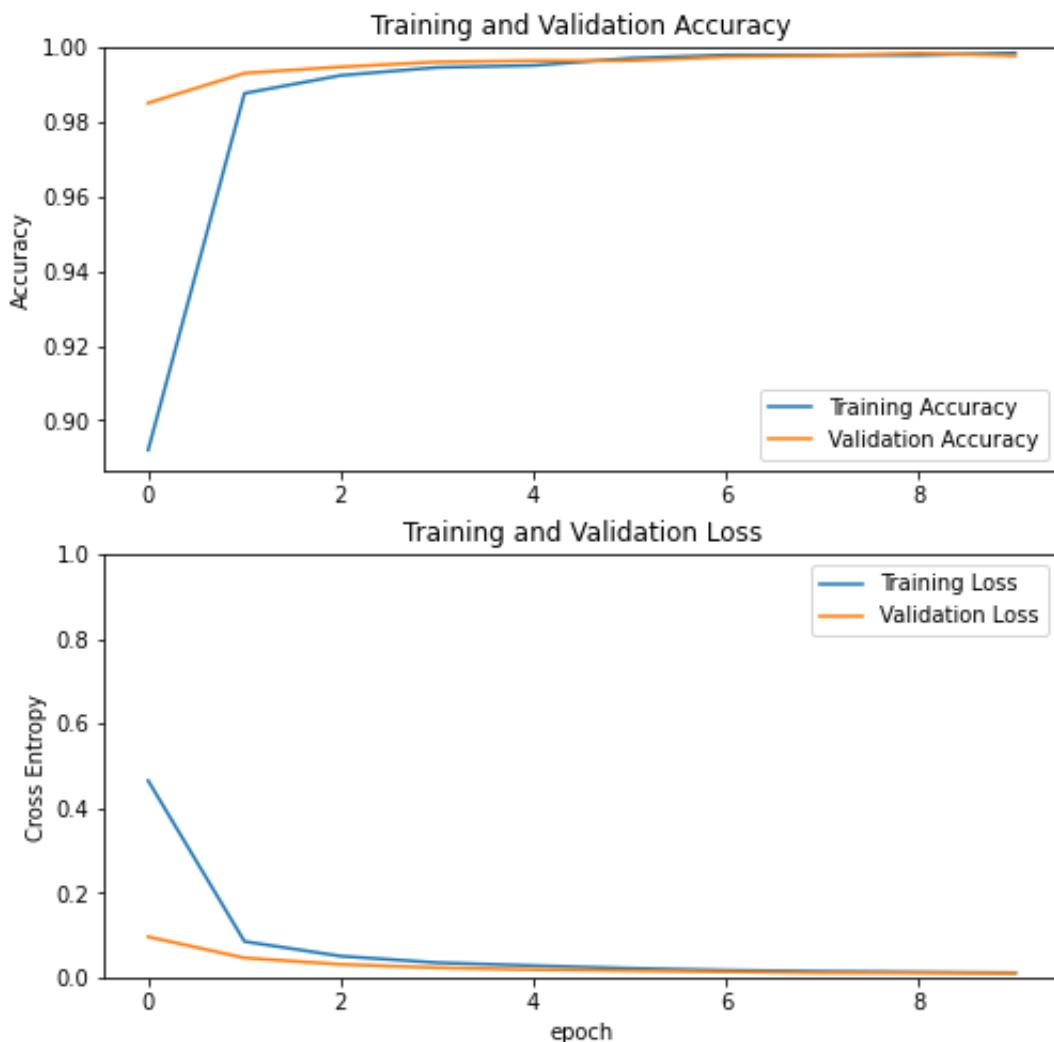
```
In [18]:
```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),1])
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.ylabel('Cross Entropy')
plt.ylim([0,1.0])
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()
```



```
In [19]: plt.figure(figsize=(10, 100))
for images, labels in test_ds.take(1):
    prediction = model.predict(images, batch_size=32)
    for i in range(9):
        ax = plt.subplot(9, 1, i+1)
        pred = np.argmax(prediction[i])
        plt.imshow(images[i].numpy().astype('uint8'))
        plt.title(f'Predito: {class_names[pred]} - Real: {class_nam
plt.axis('off')
```

1/1 [=====] - 1s 1s/step

Predito: Cauliflower - Real: Cauliflower



```
In [20]: #Visualize the result
results = model.evaluate(test_ds, verbose=0)
```

```
In [23]: print(" Test Loss: {:.5f}".format(results[0]))
print(" Accuracy on the test set: {:.2f}%".format(results[1] * 100))
```

Test Loss: 0.00997
Accuracy on the test set: 99.80%

In [24]:

```
test_ds_unbatch = tf.keras.utils.image_dataset_from_directory(  
    test_dir,  
    seed=None,  
    shuffle=False,  
    image_size=(image_height, image_width),  
    batch_size=image_height,  
)  
images = list(test_ds_unbatch.map(lambda x, y: x))  
labels = list(test_ds_unbatch.map(lambda x, y: y))
```

Found 3000 files belonging to 15 classes.

```
In [25]: y_test = np.concatenate([y for x, y in test_ds_unbatch], axis=0)  
prediction = model.predict(test_ds_unbatch)  
pred = np.argmax(prediction, axis=1)  
print(pred)  
from sklearn.metrics import accuracy_score  
  
accuracy_score(pred, y_test)
```

14/14 [=====] - 5s 257ms/step
[0 0 0 ... 14 14 14]

Out [25]: 0.998

```
In [27]: #save model  
model.save('/content/drive/MyDrive/Colab Notebooks/Vagitable.h5')
```

In []: